

AI NEXT

Rustomjee AI Champions Initiative

Program Playbook

A comprehensive design document for structuring, launching, and sustaining the AI Champions program across 30 team members, 6 teams, and 12 months.

Prepared for: Keystone Realtors Ltd. (Rustomjee)
Information Systems | AI Innovation

Executive Summary

AI Next is a 12-month program to embed AI capability directly into every operating department of Rustonjee. Thirty champions, drawn from across departments, work in six teams of five under dedicated mentors, in a purpose-built lab space, on projects prioritized by quantifiable business impact.

This document is the full operating design for the program. It covers physical and digital infrastructure, team composition, the project intake and prioritization model, the technology stack, the 12-month learning curriculum, measurement, governance, sustaining rituals, launch playbook, risks, and Year 2 graduation.

Two things will determine whether this program delivers or drifts: (1) formal protection of champion time from day-job pull, and (2) the discipline to kill weak projects fast rather than let them drag. Everything else in this document supports those two commitments.

How to read this document

- Sections 1 to 3 cover the foundation: space, teams, and how projects get selected.
- Sections 4 to 6 cover the working machinery: tech stack, learning curriculum, and measurement.
- Sections 7 to 9 cover the ongoing operating rhythm: governance, sustaining excitement, and the newsletter.
- Sections 10 to 12 cover the launch, risks, and what happens after Year 1.
- The final section lists the specific actions to take this week to move from plan to execution.

1. The AI Lab — Physical Space

The space should signal how you want people to work, not just where. Every design decision should reinforce two things: intentional collaboration and visible progress.

1.1 Layout principles

For a footprint of roughly 1,500 to 2,500 square feet, allocate space in five zones.

Zone 1 — Team pods (approximately 60% of space)

- Six clusters of five desks each, one per team, spaced to allow team huddles without disturbing neighbors.
- Each pod has one shared 32-inch monitor for pair-programming and demos.
- Whiteboard on each pod's adjacent wall — whiteboard paint is inexpensive and transformative.
- Each pod picks its own name (Team Prometheus, Team Iris, and so on) — identity matters.

Zone 2 — The War Room (approximately 15%)

- One 65-inch or larger screen, U-shaped seating for 10 to 12 people.
- Used for Demo Days, director Friday interactions, external speakers, and team reviews.
- Soundproofed if possible — sensitive project reviews and stakeholder conversations happen here.

Zone 3 — Focus corners (approximately 10%)

- Three to four single-person nooks with high partition walls for deep work.
- A physical Do-Not-Disturb flag system — champions doing detailed prompt engineering or debugging need silence.

Zone 4 — The Wall (approximately 10%)

This is the single most important design element in the lab. One long wall is dedicated entirely to making the work visible.

- A physical Kanban board showing every project by team and status: Discovery, Build, Pilot, Live.
- A Before / After board — problem statement on the left, screenshot of the shipped tool on the right, quantified benefit at the bottom. Projects graduate onto this wall when adopted.
- Milestone photos and champion badges displayed by name.
- This wall becomes your recruitment tool for Cohort 2 and your artifact when presenting to the board.

Zone 5 — Café and F&B area (approximately 5%)

- Small counter, coffee machine, fridge. Where Friday F&B lives and where cross-team conversations actually happen.

- Do not skip this even if space is tight — the informal conversations here produce more cross-pollination than any structured event.

1.2 Details that matter more than budgeted line items

- Warm lighting, not fluorescent. Champions spend 4+ hours a day here.
- Fast Wi-Fi with a dedicated SSID for the lab, separate from corporate network — champions will hit rate limits and need bandwidth for large model calls.
- One physical noticeboard for external opportunities — conference invites, courses, books being read.
- A shoe rack near the entrance — small cultural signal that this is a different kind of workspace.
- Name the lab something with intent. Not “AI Lab Room 4B.” Something that fits the Rustonjee brand voice — “The Next Room”, “Studio One”, or similar.

2. Digital and Technology Infrastructure

This is where most corporate AI programs quietly die. Lock these components before Day 1.

2.1 Core development environment

- Python 3.11 or later, environments managed with uv — faster than pip or poetry, less painful for non-developer champions.
- Gradio as the default framework for AI-focused POCs — Python-first, minimal boilerplate, ships to internal users in hours rather than weeks.
- FastAPI with a lightweight Next.js or React frontend when a tool graduates to production or requires custom UX beyond what Gradio can express cleanly.
- GitHub Enterprise with a rustomjee-ai organization and private repositories per team.
- VS Code standardized across all champions with a shared .devcontainer configuration so every environment is identical.

2.2 AI model access

- Gemini 2.5 Pro and Flash via Vertex AI as the default — already familiar to the team, good India-region data residency, generous context window.
- Claude Opus 4.7 and Sonnet 4.6 via API for tasks requiring stronger reasoning, long-document analysis, and coding assistance. Having both Gemini and Claude available is deliberate — they excel at different problems.
- Gemini Gems for no-code champions who need to ship internal-only tools quickly.
- NotebookLM for research-heavy projects such as regulatory analysis, tender review, and policy comparison.
- GitHub Copilot enterprise licenses for coding productivity across all champions.

2.3 The starter repository

Set up a shared repository named rustomjee-ai-starter on Day 1 containing reusable building blocks that every champion clones as their first act.

- Authentication wrapper integrating with corporate SSO.
- Rustomjee-branded UI theme for both Gradio and Next.js apps.
- Gemini and Claude API helper module with logging, retry logic, and cost tracking built in.
- Excel and PDF parsing utilities for common Rustomjee document formats.
- A sample Gradio application that reads a file, calls an LLM, and writes an Excel output — the “hello world” of the ecosystem.
- A sample FastAPI + Next.js scaffold for when a project outgrows Gradio.

This one repository reduces time-to-first-shipping-thing from about three weeks to about three days. Investment here compounds across every project.

2.4 Data access — the real bottleneck

Champions burn weeks trying to get access to the data they need. Pre-negotiate with Information Systems and each department head before champions arrive.

- Publish a data access matrix on Day 1: what data exists, in what system, in what form, at what classification tier, and who owns it.
- Extend the AI Governance Policy so a champion can look up their data need and see the exact approval path.
- Provide a named data owner for each system with a weekly escalation path to the CIO.

2.5 What to explicitly avoid in Year 1

- Fine-tuning custom models, training from scratch, or building custom vector search. If a champion says “I want to fine-tune,” the answer is “prompt better first, then reconsider in six months.”
- Unnecessary framework complexity introduced for its own sake. If Gradio and a Gemini call ship the tool, that beats an elegant multi-agent architecture that ships in Q3.
- Ninety percent of enterprise AI value comes from good prompting, clean data, and thoughtful UX. Year 1 should protect that focus.

3. Team Composition and Mentor Roles

3.1 Composition per pod

Six teams of five is the right structure. Four is too few for redundancy; six dilutes ownership. Within each team of five, aim for these role archetypes:

- One Tech Lead champion — strongest Python and AI skills, owns technical decisions.
- One Domain Lead champion — deepest understanding of the business problem being solved.
- One Data champion — comfortable with Excel, SQL, and data cleanup. Real-estate data is messy; this role prevents the team from stalling.
- One UX and Communication champion — good at user interviews, presentations, and written communication.
- One Wildcard — junior, high-potential, cross-department curiosity. Grows into whichever role the team needs most.

Perfect archetype coverage across all six teams is unlikely. Assign the roles you can, and let others grow into their role over the first two months.

3.2 Mixed teams vs departmental teams

There are two ways to compose teams: cluster champions by department (all Sales together), or mix by department and group by problem type (all document-intelligence work in one team).

Recommendation: mix by department, group by problem archetype.

- Departmental clustering produces sales tools that only sales can use. Mixed teams building for sales build reusable patterns other departments can adopt.
- Cross-pollination is where compound learning happens over 12 months.
- Ensure each team has at least one strong technical champion and one strong domain champion — five juniors together will stall.

3.3 The mentor role

Mentors are what turn a training program into a leadership pipeline. Be explicit about the role or it dilutes into occasional advice-giving.

A mentor is expected to:

- Meet the team twice weekly for 30 minutes each.
- Unblock the team on data access, stakeholder pushback, and tool budgets.
- Review technical decisions and business framing weekly.
- Coach on business framing, stakeholder management, and presentation.
- Protect the champion's time from their day job — escalating to directors when the 20 to 30 percent commitment is violated.

- Report honest status to the program lead, including bad news.

3.4 Mentor selection

Pick six mentors from mixed backgrounds. A useful mix:

- One from Information Systems (CIO or senior architect).
- One from Operations or a similar cross-functional role.
- One from Finance or Commercial.
- One from a Director's office (a strong senior manager).
- One from Sales or Marketing.
- One external advisor, if budget allows — an industry practitioner brings outside perspective.

Give every mentor a one-page role description in writing. Give them a small quarterly budget to take their team out for a meal — relationships form over shared food, not calendar invites.

4. Project Pipeline and Prioritization

Your Project Criteria and Prioritization Form is the base. Layer a defensible scoring model on top so selection is transparent, not political.

4.1 Scoring dimensions

Score every submitted project on the following dimensions, each on a 1-to-5 scale, weighted as shown.

Dimension	Weight	What you are measuring
Quantifiable business impact	25%	Rupees saved, hours saved, revenue enabled, or risk reduced — must resolve to a specific number
Strategic alignment	15%	Ties to a director-level priority: redevelopment, sales velocity, cost control, compliance
Feasibility with current stack	15%	Buildable with Gemini or Claude plus Gradio and available data, or requires net-new infrastructure
Data readiness	15%	Does clean, accessible data exist? This is the silent killer — score honestly
Time to first value	10%	Can a pilot ship in eight weeks or less? Long projects lose sponsors
Adoption likelihood	10%	Will the department actually use it, or is this an idea from a head who will not champion it?
Reusability	10%	Does it produce building blocks other departments can reuse?

4.2 Two non-negotiable rules

- **Sponsor commitment in writing.** A department head must sign a one-page commitment to attend biweekly reviews, provide data, and do UAT. No sponsor equals no project, regardless of how good the idea sounds.
- **Real user named at intake.** A project needs a specific user identified by name before it starts. Not “the sales team” but “Rakesh from Kandivali East sales.” Real users kill vague requirements early.

4.3 Kill criteria — defined at intake

At intake, write down what would make you stop this project. Examples: user testing shows below 60 percent satisfaction; data quality below a defined threshold; sponsor disengages for more than three weeks; cost per query exceeds a defined limit.

Real-estate AI projects fail on data quality far more often than on model choice. Explicit kill criteria prevent zombie projects that consume resources without shipping value.

4.4 Milestone definitions

You said you will celebrate every milestone, not just completion. That instinct is correct — but the milestones need to be defined clearly enough to earn celebration.

- Problem Validated — users interviewed, problem confirmed, kill criteria set.
- POC Working — the core AI pipeline produces the intended output in a lab setting.
- Pilot with 5 Users — real users using the tool for real work, feedback captured.
- Rolled Out to Department — tool available to the full target user base.
- Measured Business Impact — quantified benefit reported and validated by sponsor.

Each milestone earns a small ritual: a Slack callout, a physical pin on the champion's badge lanyard, a shout-out in the newsletter. Sustains a 12-month program in a way that big-bang launches cannot.

5. Standardization Choices

Standardization matters more than optimal choice at this stage. Uniform tools let champions help each other, let mentors review any team's code, and let learning content apply universally.

5.1 Model access mix

- Gemini 2.5 Pro / Flash via Vertex AI — default choice.
- Gemini Gems — no-code path for internal tools.
- NotebookLM — research-heavy work.
- GitHub Copilot — coding productivity for everyone.

6. 12-Month Learning Curriculum

Two parallel tracks. The Tech Track builds the ability to ship AI tools. The Business Track builds the ability to lead adoption. Both are heavy on doing, light on lectures.

6.1 Tech Track — weekly, 90 minutes

First hour teaches; last 30 minutes is hands-on with a task the champion submits by end of week.

Months 1 to 3 — Foundations

- Week 1: What LLMs actually are, how APIs work, hello-world Gemini and Claude calls.
- Week 2: Prompt engineering — clarity, examples, structure, XML tags.
- Week 3: Python for AI — the 20 percent you need for 80 percent of the work.
- Week 4: Gradio basics — building a working file-upload tool.
- Week 5: Working with PDFs, Excel files, and images.
- Week 6: Structured outputs — JSON mode, schemas, validation.
- Week 7: Multi-turn and stateful applications.
- Week 8: Prompt evaluation — how do you know your prompt is any good?
- Week 9: Error handling, cost management, rate limits.
- Week 10: Git, branches, pull requests — practical development workflow.
- Week 11: Deploying a Gradio app to internal users.
- Week 12: Mid-cohort Demo Day.

Months 4 to 6 — Intermediate

- RAG and vector search — when it helps, when it is overkill.
- Function calling and tool use.
- Working with images — OCR, drawings, floor plans.
- Basic agents — planning, tool use, memory.
- Evaluation frameworks — building your own eval set.
- Guardrails, safety, prompt injection defense.
- Handling multilingual data — English, Hindi, Marathi.
- Advanced Gradio patterns and FastAPI integration — state, custom components, backend patterns.

Months 7 to 9 — Applied

- Deep dives on real-estate AI use cases (bring in external speakers).
- Advanced document intelligence.
- Production considerations — logging, monitoring, cost tracking.

- Data engineering basics — ETL, data quality, pipelines.
- Working with SAP, Salesforce, and internal CRM data.
- Model comparison and selection.

Months 10 to 12 — Consolidation

- Scaling POCs to production — the FastAPI + Next.js path.
- Change management and adoption.
- Presentation and storytelling.
- Peer teaching — champions teach each other.
- Building Cohort 2's curriculum.

6.2 Business Track — fortnightly

Often best held during the Friday director interaction slot, alternating with other Friday content.

- Writing a business case that gets funded.
- Quantifying ROI for AI projects (harder than it sounds).
- Stakeholder mapping and influence.
- User interviews — how to ask real questions, not confirmation-seeking ones.
- Presenting to leadership — the five-slide deck.
- Change management basics.
- Building coalitions across departments.
- Vendor evaluation.
- AI ethics and governance in a real-estate context.
- Reading the market — what competitors are doing.
- Personal brand — LinkedIn and thought leadership.
- Career pathing in an AI-augmented organization.

6.3 External speakers

Aim for one external voice per month. A good mix over the year includes:

- A real-estate technology founder.
- Someone from the Google, Microsoft, or AWS India AI team.
- An academic — IIT Bombay's AI lab is nearby and well-regarded.
- Someone from a non-real-estate industry showing transferable ideas.

Record every session. Make attendance mandatory but respected — no phones out. These sessions signal ambition and build a network for champions that will outlast the program.

7. Measurement Framework

Three tiers of measurement. Without them, the program will not survive its first budget review.

7.1 Tier 1 — Project-level metrics

Tracked from Day 1 of every project.

- Business impact: a single number — hours saved, rupees saved, revenue enabled.
- Users: how many people use the tool weekly.
- Adoption trajectory: users over time.
- Quality: user satisfaction (a simple thumbs-up or thumbs-down in the app).
- Cost to build: person-hours plus tool costs.
- Cost to run: monthly API and infrastructure spend.

7.2 Tier 2 — Program-level metrics

Reported quarterly to directors.

- Total projects: proposed, active, killed, shipped, adopted.
- Total business impact: sum of Tier 1 numbers.
- Champion retention — a leading indicator. If champions leave, the program is failing.
- Champion engagement — hours logged, sessions attended.
- Champion NPS — anonymously, quarterly.
- Departments actively participating.
- Cross-department collaboration count.

7.3 Tier 3 — Strategic metrics

Annual, for the board.

- Percentage of Rustomjee employees using at least one AI tool built by the program.
- Reduction in cycle time for target processes.
- Competitive positioning versus peers, survey-based.
- Talent brand impact — is Rustomjee attracting AI talent now?
- Program ROI: total business impact divided by total program cost.

7.4 Reporting rhythm

- Weekly: Kanban wall physically updated in the lab.
- Fortnightly: newsletter includes headline metrics.
- Monthly: dashboard shared with all mentors.
- Quarterly: one-page dashboard to the three directors.

- Annually: full board presentation.

8. Governance and Guardrails

8.1 The Champion Charter

Extend the existing AI Governance Policy with a one-page Champion Charter that every champion signs. Their manager signs it too. Contents:

- What data can and cannot be used (linked to data classification tiers).
- Requirement to log every AI tool built in the central project registry.
- Requirement to consult Information Systems before deploying anything that touches customer or employee data.
- Prohibition on using client personal data in prompts without documented approval.
- IP: everything built during the program is Rustomjee IP.
- Confidentiality: internal tools stay internal.
- Time commitment: 20 to 30 percent of role formally, signed by the champion's manager as well as the champion.

8.2 Central AI project registry

A simple internal web tool where every AI project is registered with owner, data used, users, purpose, and launch date.

This is the compliance shield and the reusability map. When a new champion wants to build something similar, they find prior work here rather than starting from zero.

8.3 The Weekly AI Council

Thirty minutes, weekly. Members:

- Program lead.
- CIO.
- One director, on rotation.
- One champion, on rotation.

Reviews new project proposals, escalations, and stuck projects. Decisions documented in the newsletter for transparency.

9. Sustaining Excitement

Excitement is easy at launch. The middle six months are where programs die. Structured rituals and progressive recognition prevent that.

9.1 Weekly rituals

- Monday standup per team (15 minutes, mentor-led).
- Wednesday cross-team demo (30 minutes, one team shows current work).
- Friday: F&B in the lab, director interaction, and either an external speaker or a Business Track session.

9.2 Monthly rituals

- Champion of the Month — peer-voted, not management-picked.
- One external speaker.
- Monthly newsletter recap by a director.

9.3 Quarterly rituals

- Demo Day — every team presents, directors judge, small prizes awarded.
- Champion quarterly review — feedback in both directions.
- Program metrics review with directors.
- Champion offsite or team dinner.

9.4 Rewards ladder

- Milestone pins, physical and wearable: Problem Validated, First Pilot, First 100 Users, First 1000 Users, Deployed.
- Champion of the Month: modest cash bonus (5,000 to 10,000 rupees), newsletter feature, LinkedIn shoutout from directors.
- Champion of the Quarter: larger reward, external conference sponsorship.
- Champion of the Year: significant recognition, promotion consideration, LinkedIn profile written by CEO.

9.5 Learning stipend

10,000 rupees per champion per year for books, courses, and conferences. Small money, enormous signal. This tells champions the organization is investing in them personally, not just extracting output.

9.6 Conference sponsorship and LinkedIn strategy

- Send 2 to 3 champions per quarter to external AI events — Google Next India, Nasscom AI, industry-specific conferences. They return with ideas and act as evangelists.
- Every shipped milestone gets a coordinated LinkedIn post from the champion, directors, and Rustomjee corporate. Champions become known externally.
- External visibility is both retention insurance and recruiting fuel.

10. The Newsletter

Fortnightly, five-minute read maximum, distributed to every Rustomjee employee — not just champions. This builds organization-wide AI literacy and creates demand for the champions' work.

10.1 Format

- Header: issue number, date, one-line tagline from a director.
- Wins: two to three milestones hit that fortnight, one line each with the number.
- Project Spotlight: one project, 400 words maximum, co-written by a champion. Problem to approach to result to next step. Include a screenshot.
- Champion Q&A: five questions with one champion. Rotate so everyone gets featured over the year.
- Tool of the Fortnight: one AI tool the entire company can try. Democratizes AI beyond the lab.
- Industry Radar: two to three links from real-estate AI news, with a one-line "why we care" note each.
- Upcoming: speakers, demos, milestones next fortnight.
- Ask the Directors: one anonymous question submitted by a champion, answered by a director.

10.2 Editorial model

- Rotate co-authorship — each edition has a different champion as co-editor whose name goes on the byline.
- The program lead is the executive editor and quality gate.
- Design consistency handled by corporate communications or brand team.
- First edition ships in Week 4 of the program.

11. Launch Playbook — First 30 Days

11.1 Day 1 — Kickoff Event

Two-hour launch event with the following flow:

- All three directors present.
- CEO or MD opens (5 minutes) — why this matters for Rustomjee.
- Each director speaks (5 minutes each) — what they personally hope to see.
- Program lead walks through structure, timelines, and expectations (15 minutes).
- Champions introduced by name and department (30 minutes).
- Team assignments revealed (10 minutes).
- Mentor introductions (10 minutes).
- T-shirts distributed.
- Group photo — used in every subsequent newsletter and LinkedIn post.
- Reception with F&B, mixing across teams.

11.2 Week 1

- Technical infrastructure ready — every champion has GitHub access and their environment set up.
- Data access matrix published.
- Champion Charter signed by every champion and their manager.
- First Tech Track webinar — Hello World with Gemini and Claude.
- Project intake opened for department heads.
- Team channel created in corporate messaging platform — #ai-next-champions.

11.3 Week 2

- Team pods finalized in the physical space.
- Kanban wall set up.
- Each team picks its name and identity.
- Second Tech Track webinar — Prompt Engineering.
- First one-on-one meetings between each mentor and each champion.

11.4 Week 3

- First project proposals submitted through the scoring model.
- AI Council meets to select the first six projects — one per team.
- Sponsors and real users identified per project.

- Third Tech Track webinar — Python for AI.
- First Business Track session.

11.5 Week 4

- Every team has a project, a sponsor, a named user, and defined kill criteria.
- Fourth Tech Track webinar — Gradio basics.
- First director Friday interaction.
- First newsletter goes out.

11.6 Weeks 5 through 12

- Weekly rhythm established: standups, cross-team demos, webinars, Friday sessions.
- First POCs shipping by Week 8.
- First formal Demo Day at Week 12.
- First Champion of the Month announced end of Month 3.

12. Risks and Mitigations

The five most likely reasons this program could fail, ranked, with mitigations.

Risk 1: Day-job conflict

Champions get pulled into fires by their department heads. The program becomes second priority.

Mitigation. Written 20 to 30 percent time commitment signed by the manager. Weekly check-in on protected time. Direct escalation to directors when violated. This is the single largest risk and the single most important mitigation.

Risk 2: Data access blockers

Champions cannot get the data they need. Projects stall at around 40 percent completion.

Mitigation. Data access matrix published Day 1. Named data owners per system. Weekly escalation path to CIO.

Risk 3: Departmental sponsor drift

Sponsors lose interest around Month 2.

Mitigation. Written sponsor commitment. Biweekly review attendance. Kill projects whose sponsors disengage — saves the champion's time for a better project.

Risk 4: The AI theater trap

Tools built, adoption never happens, benefits are theoretical.

Mitigation. Real user named at intake. Kill criteria include adoption thresholds. Measurement Tier 1 includes actual usage, not just deployment.

Risk 5: Best champions leave

They receive external offers based on the skills they built during the program.

Mitigation. Clear career path from champion to senior champion to mentor to AI lead. Retention bonus at 12 months. LinkedIn visibility so growth is public. Salary review at 6 months for standouts.

13. Graduation and Year 2

At Month 12, every champion should have a defined next step. The absence of a career path is what kills strong talent.

13.1 Champion pathways

- Senior Champion — stays in the program, mentors a Cohort 2 team.
- AI Lead in Department — returns to department full-time as the AI point person, with a title change.
- AI Core Team — joins a permanent central AI function, likely reporting to the CIO.
- Alumni Network — if a champion moves on, they remain in the alumni network with quarterly events and first refusal on new opportunities.

13.2 Cohort 2 launch

- Cohort 2 launches Month 13, informed by everything learned in Cohort 1.
- Cohort 1 champions run roughly half the training sessions — teaching accelerates their own mastery.
- The program scales without the program lead personally doing more work each year.
- The Wall of shipped projects becomes the recruitment material for Cohort 2.

14. Immediate Next Steps

The specific actions to take this week to move from plan to execution.

- Secure written buy-in from the three directors on time protection for champions — this is the single biggest risk mitigation.
- Draft the Champion Charter and get HR sign-off before champion selection is finalized.
- Lock the physical space and start infrastructure ordering — Wi-Fi, monitors, and whiteboards have long lead times.
- Circulate the scoring model to department heads submitting projects so intake begins with the right framing.
- Identify and confirm the six mentors — this is a limiting factor for the launch date.
- Book the launch event date and block the directors' calendars.
- Set up the rustomjee-ai-starter repository skeleton with the Gradio and FastAPI templates.
- Draft the first newsletter template and get design sign-off from the brand team.
- Define the initial data access matrix with the CIO — what data exists, who owns it, at what classification tier.
- Draft the sponsor commitment template so department heads can sign at project intake.

Closing Note

The ideas in this document will not make or break the program. Execution discipline on two things will: protecting champion time formally and killing weak projects fast.

Do those two well, and everything else in this document — the T-shirts, the F&B, the Friday director sessions, the newsletter, the pins — becomes sustaining ritual rather than launch theater.

The 30 champions are the greatest asset. The lab is the visible artifact. The projects are the near-term ROI. But the long-term win is that Rustomjee builds a culture where AI is not a project, it is how work happens.